

Evanesca: Measuring DeFi Value Extraction at Scale

Jeongmin Kim

DeFi protocols have suffered over \$10 billion in losses between 2020 and 2025, yet most incidents are studied as isolated exploits. We observe that many attacks share recurring structure at the level of financial actions such as swap, borrow, and repay, suggesting that value extraction can be measured systematically. We present Evanesca, a measurement pipeline that reconstructs Semantic Financial Graphs (SFGs) from on-chain activity and evaluates reusable constraints expressed in a domain-specific language. Applying nine constraints in 107 lines to 14,187,803 Ethereum transactions (2020–2024) and 40 confirmed incidents spanning five chains, Evanesca reconstructs every incident without incident-specific signatures. The same rule set flags 1,418 operational instances: 956 baseline arbitrage and 462 non-baseline instances spanning reward-distribution manipulation, price manipulation, yield-bearing token deviations, and a long tail of mixed signals. Source-code investigation of flagged instances uncovered 219 previously unreported implementation defects across two categories: derivative pricing errors and reward-distribution bias. Replayable evidence traces let Evanesca serve as a reusable starting point for longitudinal studies of DeFi risk.

Report scope. This technical report documents the Evanesca research prototype, its public artifact boundary, and the IMC-version measurement snapshot. It is written for readers who want to inspect the method and reproduce representative evidence, not as a production detector specification. Scale-out numbers follow the IMC snapshot; implementation notes describe the open-source release state.

Additional Key Words and Phrases: Blockchain, Ethereum, Smart contract analysis, MEV, DeFi security

1 Introduction

Decentralized Finance (DeFi) suffered over \$1.6 billion in losses in Q1 2025 alone [10], part of a multi-year cumulative loss trend across the ecosystem, yet incident response and academic studies still treat most attacks as isolated exploits. We observe that many incidents share a common structure at the level of *financial actions* such as swap, borrow, and repay: the “borrow–manipulate–profit” sequence behind the bZx attack [14, 15] recurs across different protocols, chains, and years [24]. This repetitiveness suggests that value extraction can be measured systematically rather than investigated case by case.

However, existing approaches do not exploit this recurring structure. First, current DeFi attack detection systems do not generalize to unseen patterns: runtime monitors [8, 13] are configured per deployment, while attack-hunting frameworks such as DeFiRanger, BlockEye, DeFort, and DeFiScope [18, 20–22] bind detection to predefined attack templates—resulting in low coverage of incidents at the protocol layer, such as oracle manipulation and permissionless cross-contract interactions [24]. Second, none of these systems produce replayable evidence traces, and many report only aggregate statistics that cannot be independently verified. They emit alerts but do not compute financial values at each step, so analysts cannot triage variants or reproduce findings across studies [8, 13, 21].

Beyond these technical limitations lies a more fundamental challenge. Value extraction in DeFi does not divide neatly into attacks and normal behavior. Maximal Extractable Value (MEV) studies [1, 7, 19] show that significant value can be extracted by reordering transactions within a block without a discrete software vulnerability [24], blurring the boundary between “attack” and “strategy.” We therefore frame this as a measurement problem rather than intent classification.

Our measurement framing requires protocol-independent invariants over typed financial actions, unlike prior frameworks that bind detection to fixed attack templates. The underlying insight is that the same service-level invariants recur across DeFi protocols, allowing a small rule set over typed actions such as swap, borrow, and repay to capture both benign baseline arbitrage and more suspicious non-baseline patterns. Accordingly, Evanesca is a measurement pipeline of four composable stages, each defined by its input and output. *Semantic recovery* takes protocol-specific smart-contract events as input and emits uniform financial actions (swap, borrow, deposit,




Fig. 1. Measurement pipeline. Evanesca performs semantic recovery, constructs USD-normalized SFGs, evaluates DSL constraints, and outputs replayable evidence traces.

etc.). *SFG construction* assembles these actions into a *Semantic Financial Graph* whose edges carry typed actions annotated with USD-normalized values. *Constraint evaluation* then applies nine protocol-independent DSL rules¹ that encode invariants common to entire DeFi service classes (e.g., value preservation across any swap, collateral requirements for any borrow). *Evidence extraction* materializes each constraint match as a replayable trace. Each match is a *flagged instance*, a measurement signal indicating a pattern worth investigating, not a verdict that the transaction was harmful.

In summary, we make the following contributions:

- **Multi-chain measurement at scale.** Using only general-purpose constraints (nine DSL rules in 107 lines, without incident-specific logic), we reconstruct all 40 confirmed incidents across five chains and flag 1,418 operational instances across five pattern categories over 14.2 M Ethereum transactions (2020–2024).
- **Reproducible methodology.** Every flagged instance is backed by a replayable evidence trace that records the full action sequence and profit path, enabling independent verification.
- **Previously unreported defects.** Source-code investigation of flagged instances uncovers 219 implementation defects across two categories: derivative pricing errors and reward-distribution bias.

2 Threat Model and Background

Threat model. Public blockchains maintain an append-only ledger of transactions readable by anyone. Smart contracts execute deterministically, making all financial activity publicly observable and reproducible. We study value-extraction patterns observable on this ledger: sequences of financial actions that violate invariants common to a DeFi service class, such as pool invariants for swaps, collateral requirements for borrows, or exchange-rate consistency for deposit/withdraw cycles, or that extract rewards and fees beyond intended use. A pattern is flagged when a constraint fires on the evidence trace. The flag is a measurement signal, not an intent label. We do *not* attempt real-world attribution. We measure patterns at the level of addresses and transaction sequences.

DeFi protocol primitives. Decentralized Finance (DeFi) protocols, autonomous financial services implemented as smart contracts, compose a small set of financial actions: Automated Market Makers (AMMs) price token *swaps* via pool invariants (e.g., $x \cdot y = k$ where x, y are token reserves); lending protocols let users *deposit* collateral, *borrow*, *repay*, and *withdraw*; *flash loans* provide uncollateralized capital that must be repaid within the same transaction; and *yield-bearing tokens* (e.g., Compound’s cTokens, Aave’s aTokens) are receipt tokens representing deposit positions whose exchange rate accrues over time. Price *oracles* supply external market data to contracts. Manipulating an oracle corrupts downstream pricing decisions. Every computation costs *gas*, priced in the native currency, and users pay gas fees proportional to transaction complexity. These primitives are the building blocks of both legitimate DeFi activity and the patterns we measure.

Value-Extraction Patterns. We define a *value-extraction pattern* as a repeatable sequence of *financial actions* executed to realize profit with minimal exposure. This definition is intentionally broad: routine arbitrage also

¹We use *rule* and *constraint* interchangeably throughout this report. The collection of all nine forms our *rule set*.

fits and is retained as an explicit baseline in our evaluation. The distinction between baseline and non-baseline patterns is made during analysis: baseline patterns profit from existing price discrepancies across markets, whereas non-baseline patterns require the actor to alter protocol state (e.g., inflating a token price in a low-liquidity pool) to create the profit opportunity.

Measurement objective. Given raw on-chain activity, we reconstruct the sequence of actions underlying each pattern, quantify counts and conservative impact, and output evidence artifacts that an analyst can inspect and reproduce. We do not classify adversarial intent, estimate total DeFi losses, or build a real-time monitoring system. Our focus is retrospective, evidence-based measurement.

3 Measurement Methodology

We operationalize the measurement objective as three research questions, then build a pipeline that answers them.

- **(RQ1)** Can general-purpose constraints reconstruct known DeFi incidents?
- **(RQ2)** What patterns emerge at scale?
- **(RQ3)** Do flagged patterns correspond to real implementation defects?

This section answers RQ1 and develops the pipeline shared by all three. Section 6 applies the same pipeline at scale for RQ2 and RQ3.

To answer these questions, we translate raw on-chain events into structured sequences of financial actions such as swaps, deposits, and borrows, with their associated values. Unlike contract-level analysis [6, 12, 17] or runtime monitoring [8, 13, 16], our approach operates on semantic financial actions that generalize across protocols and chains. Three challenges motivate our pipeline design.

(C1) Semantic gap: Protocols implement the same operation using different event schemas (e.g., `Mint` on Compound vs. `Deposit` on Aave); contract-specific heuristics therefore break across forks and upgrades. **(C2) Temporal composition:** Individual actions such as a swap or a borrow may appear normal in isolation. The extraction becomes visible only when the full ordered sequence within a transaction is analyzed together, often spanning multiple protocols. **(C3) Reproducibility:** Existing pipelines, whether rule-based monitors [8, 13] or template-based frameworks [21], typically emit alerts without the per-step computation needed for independent replay and triage.

Pipeline overview. Our pipeline processes on-chain data through four stages (Figure 1): semantic recovery, SFG construction, constraint evaluation, and evidence extraction. Each constraint match produces a *flagged instance*, which bundles the evidence trace from the SFG, the violated constraint, and per-entity PnL.

Semantic recovery. Protocol-specific adapters decode each transaction into five uniform action primitives: swap, deposit, withdraw, borrow, and repay. We cover 40+ protocol families spanning DEX, lending, flash-loan, and bridge services (Section 5). Flash loans are emitted as a within-transaction borrow+repay pair. Because flash loans require no upfront capital, connecting borrowed funds to downstream profit is essential for identifying patterns that extract value with no upfront capital. Unlike contract-specific signatures, our adapters produce uniform output across protocol families, addressing C1.

SFG construction. We model each transaction as a directed graph $G = (V, E)$. Vertices represent entities such as accounts, liquidity pools, and lending markets, annotated with their protocol role. Edges represent typed financial actions, recording input/output assets, USD-normalized values, and the originating protocol. We convert token amounts to USD using the CoinGecko price API [4] at each transaction’s block timestamp. Edges are ordered by transaction index, ensuring deterministic evaluation. This representation abstracts away protocol-specific event schemas while preserving the ordering and cross-protocol connectivity of financial actions. As a result, the same




Fig. 2. Representative price manipulation pattern (bZx), the RQ1 walkthrough. Semantic recovery connects the manipulation step to profit across protocols.

multi-step extraction pattern remains visible even when implemented by different contracts, addressing C1 and enabling C2.

More formally, an SFG can be viewed as $G = (V, E, \phi, \tau)$, where ϕ maps each edge to its recovered action semantics (swap, deposit, withdraw, borrow, or repay) and τ gives the strict temporal order induced by the transaction trace. Each edge carries the action type, service class, concrete protocol, input and output token amounts, USD-normalized values, and trace metadata needed for replay. This matters for measurement because the stable unit of comparison is often the financial action rather than the low-level event: a swap may appear as different event names across Uniswap, Curve, Balancer, or a fork, but it exposes the same value-preservation invariant once lifted into an SFG edge.

Constraint evaluation. We express value-extraction patterns as reusable semantic rules over SFG edges. For each edge e whose input has nonzero USD value, we compute:

$$\rho(e) = \frac{\text{valueUSD}(\text{out}(e))}{\text{valueUSD}(\text{in}(e))} \quad (1)$$

and flag edges where $|\rho(e) - 1|$ exceeds an edge-type-specific threshold. For DEX swap edges we use $\alpha = 0.05$, a conservative bound above the 1.16% average unexpected slippage on Uniswap reported by Zhou et al. [23]. For lending and oracle-mediated edges involving stablecoin derivatives we use a tighter $\beta = 0.02$, reflecting typical stablecoin oracle noise: an empirical study of DeFi oracles [11] shows that SAI and DAI daily USD price changes stay within $\leq 1\%$ on 66–70% of days and within $\leq 5\%$ on over 98% of days, so $\beta = 0.02$ sits just above this typical noise floor. Our rule set comprises nine DSL constraints in 107 lines, spanning AMM invariants, lending and accounting checks, oracle anomalies, bridge integrity, and flash-loan patterns. The grammar and examples appear in Section 4. Protocol-specific logic is confined to the SFG construction layer. The constraints themselves are protocol-independent.

Evidence extraction. Given the flagged edges and the SFG, the pipeline produces replayable evidence traces with per-vertex profit and loss (PnL). An evidence trace $H \subseteq G$ is the subgraph induced by all flagged edges

Table 1. Confirmed incidents (RQ1): distribution by attack type and coverage by constraint family.

Attack type	Count	Constraint family	Incidents [*]
Flash loans (incl. flash-loan+price)	13	Price-deviation constraints	22 (55.0%)
Price manipulation	11	AMM constraints	22 (55.0%)
Reentrancy	7	Collateral/accounting constraints	12 (30.0%)
Oracle manipulation	4	Exchange-rate constraints	9 (22.5%)
Bridge exploits	2	Oracle integrity constraints	7 (17.5%)
Other (logic, governance, donation)	3	Bridge / flash-loan constraints	1 (2.5%) each
Total	40		

^{*}Percentages sum to >100% because one incident may trigger multiple constraint families.

within a transaction. For each vertex v in H , we compute:

$$\begin{aligned} \text{PnL}(v) = & \sum_{e \in \text{in}_H(v)} \text{valueUSD}(e) \\ & - \sum_{e \in \text{out}_H(v)} \text{valueUSD}(e) \end{aligned} \quad (2)$$

This highlights beneficiaries (positive PnL) and likely victims (negative PnL); we attribute impact across protocols, separating the protocol where the violation occurs from the entity that bears the loss. The per-step computation enables independent replay, addressing C3. The emitted artifact includes the matched constraint, the minimal action subgraph, the relevant $\rho(e)$ and PnL values, and the original transaction identifier. We use these artifacts for manual audit and for public examples, rather than relying on a transaction-level alert without intermediate computations.

Illustrative example. The bZx incident [14, 15] illustrates the pipeline end to end. The attacker borrows via flash loan, manipulates a low-liquidity pool, and extracts profit from a lending protocol that trusts the manipulated price (Figure 2). The price-inflating swap yields $\rho = 1.56$ (56% excess output relative to input) and the victim’s counterparty swap yields $\rho = 0.35$ (65% value loss), both triggering the abnormal-swap constraint. The evidence trace links these flagged edges to the attacker’s profit path.

Validation on known incidents. We curated a confirmed-incident set of 40 publicly documented incidents with available transaction traces across our five target chains; the set spans price manipulation attacks for DeFiRanger comparison plus reentrancy, governance, and bridge exploits for scope breadth (Table 3). We define successful reconstruction as meeting three criteria: (i) the evidence trace identifies the correct manipulation step, (ii) it traces a profit path to the beneficiary, and (iii) at least one general-purpose constraint fires on the value-extracting edge. Across all 40 confirmed incidents, our pipeline meets these criteria. Table 1 summarizes incident diversity. Price-deviation and AMM constraints each fire on 55.0% of incidents, followed by reentrancy and collateral/accounting at 30.0% each.

Finding (RQ1). The nine protocol-independent rules reconstruct 40/40 confirmed incidents across five chains without per-incident logic. These rules capture structural patterns that recur across incidents, which is what makes the scale-out evaluation in §6 tractable.

4 Constraint DSL

Purpose. Evanesca encodes value-extraction pattern hypotheses as *auditable constraints* over SFG action edges (Section 3). The DSL is designed to make patterns easy to review, version, and rerun across time and services, while still producing replayable evidence artifacts.

Constraint format. Each constraint declares (i) when it applies, (ii) derived variables, and (iii) a violation predicate:

```
constraint NAME {
  when: <boolean expression over edge fields>
  condition: { x: <expr>, y: <expr>, ... }
  violation: <boolean expression>
  message: "human-readable summary"
}
```

when filters the edges to which the constraint applies; condition binds intermediate expressions by name (used by violation); message is attached to flagged matches for triage. Expressions support arithmetic and boolean operators (+, -, *, /, <, <=, ==, &&, ||) and member access such as edge.Type and edge.totalInUSD. In condition bindings, || serves as a fallback operator: for example, x || 0 yields x if defined, else 0. In violation expressions, || is logical OR.

Evaluation context. The DSL is evaluated over the per-edge records emitted by semantic recovery and USD normalization. In the examples below, edge refers to the current SFG edge and exposes fields such as Type, Action, Service, totalInUSD, totalOutUSD, and other semantics derived from logs/traces.

Representative examples. The DSL constraint set used in our evaluation contains nine constraints; we show three representative patterns below. The first is the abnormal-swap rule, which computes a value ratio price_ratio (equivalently, $\rho(e) = \text{valueUSD}(\text{out})/\text{valueUSD}(\text{in})$) and adds an auxiliary trigger that fires when a large price impact coincides with transaction atomicity, suppressing noise at scale.

```
constraint PRICE_MANIPULATION {
  when: edge.Type == "DEX" &&
        edge.Action == "Swap"
  condition: {
    price_ratio:
      (edge.totalOutUSD / edge.totalInUSD),
    price_impact:
      edge.price_impact_percent || 0,
    is_atomic:
      edge.is_atomic_transaction || false
  }
  violation:
    edge.totalOutUSD > 0 &&
    edge.totalInUSD > 0 &&
    (price_ratio > 1.05 ||
     (price_impact > 10 && is_atomic))
  message:
    "Price manipulation detected via abnormal exchange rate"
}
```

Production constraints may add additional thresholding or accounting guards; we show the core rule structure for clarity.

The next example is a concrete AMM-invariant check:

```
constraint DEX_K_INVARIANT {
  when: edge.Type == "DEX" &&
        edge.Action == "Swap"
  condition: {
    k_ratio: edge.k_after / edge.k_before
  }
  violation:
    edge.k_before > 0 &&
```

```

    edge.k_after > 0 &&
    (k_ratio > 1.01 || k_ratio < 0.99)
  message:
    "D1: K-invariant violation detected in AMM"
}

```

Finally, a lending-side accounting check:

```

constraint LENDING_COLLATERALIZATION {
  when: edge.Type == "Lending" &&
    (edge.Action == "Borrow" ||
     edge.Action == "Withdraw")
  condition: {
    collateral_ratio:
      edge.collateral_value /
      edge.borrow_amount_usd,
    health_factor:
      edge.user_collateral /
      edge.user_debt,
    profit_check:
      edge.totalOutUSD > edge.totalInUSD * 1.3
  }
  violation:
    (edge.borrow_amount_usd > 0 &&
     collateral_ratio < 1.0) ||
    (edge.user_debt > 0 &&
     health_factor < 1.0) ||
    (profit_check &&
     edge.totalInUSD > 0)
  message:
    "Under-collateralized lending position detected"
}

```

These examples illustrate how patterns are encoded as explicit, replayable rules; the full library includes additional constraints for exchange-rate anomalies, flash-loan/reentrancy patterns, and bridge integrity. For readability, we insert line breaks in the examples and may shorten long message strings; the rule logic is unchanged.

5 Implementation

Evanesca is implemented as a modular TypeScript framework with four layers mirroring the pipeline stages in Section 3: (i) *semantic recovery* (event decoding and protocol-specific action mapping), (ii) *graph construction* (SFG building with strict action ordering), (iii) *constraint evaluation* (DSL matching over SFG edges), and (iv) *evidence extraction* (subgraph induction and per-vertex PnL). The implementation goal is to keep outputs reproducible from on-chain data while keeping batch execution practical at multi-year scale. We plan to release the framework as open source.

Semantic adapters. Each supported protocol family has a dedicated adapter that maps its event schema to SFG action edges. We currently ship adapters for 40+ protocol families spanning DEX (Uniswap V1/V2/V3, SushiSwap, PancakeSwap, Curve, Balancer, Bancor, KyberSwap, Synthetix, Platypus, WooFi, and others), lending (Compound, Aave, Cream Finance, Venus, Inverse, Euler, bZx, Radiant, Hundred, and others), flash-loan (dYdX, AaveV3, DODO), and bridge protocols (Qubit, Allbridge), covering constant-product AMMs, stableswap, concentrated liquidity, lending, flash-loan, and cross-chain primitives. Adding a new protocol requires implementing one adapter module; existing constraints and evidence extraction remain unchanged. Adapters emit typed action edges rather than detector verdicts. This separation is intentional: protocol-specific logic stays in semantic recovery, while constraint definitions remain service-class level and can be rerun over different datasets without rewriting exploit templates.

Ingestion and chain configuration. The collector uses a unified EVM ingestion interface with chain-specific configuration for block timing, RPC/log retrieval, and token metadata. The IMC measurement snapshot uses

Ethereum for the 14.2M-transaction scale-out study and retrieves confirmed-incident traces from BSC, Arbitrum, Avalanche, and Optimism for RQ1 coverage. The same ingestion abstraction is designed to support additional EVM-compatible networks when their traces and token metadata are available.

Scalability and batch measurement. For each transaction, Evanesca evaluates all DSL constraints on every SFG edge. The current rule set contains nine constraints in 107 lines of DSL. Each constraint includes a when clause that acts as a lightweight type filter; for example, a DEX invariant fires only on swap edges. Irrelevant constraints are then skipped without full evaluation. Performance comes from three optimizations: (i) DSL constraints are compiled to JavaScript functions at load time, avoiding repeated parsing; (ii) parsed constraint ASTs are cached with SHA-256 hashing; and (iii) token-to-USD conversions are batched before constraint evaluation to avoid redundant price lookups. The evaluator is staged: inexpensive type filters and single-edge invariants run first, and only surviving candidates trigger more expensive temporal correlation and evidence-subgraph extraction. This keeps the number of analyst-facing artifacts manageable while preserving the intermediate computations needed for reproducibility.

USD normalization. Resolved CoinGecko prices [4] are cached per (token, timestamp) pair with LRU eviction and bounded TTL to reduce API calls during batch processing. CoinGecko exposes historical prices indexed by timestamp, so the same transaction can be re-priced consistently across runs as long as the API remains accessible. Alternative historical-price sources (e.g., CoinMarketCap’s historical API [5], Chainlink on-chain price feeds [3]) could replace CoinGecko without changes to the rest of the pipeline.

Deployment and throughput. Our distributed deployment uses 3 machines (Intel Xeon E5-2640 v4 @ 2.40 GHz, 128 GiB RAM each) running 30 Docker workers, 10 per machine. Mean per-transaction processing time is 309 ms (p95 < 1.2 s; max 9.9 s), keeping multi-year measurement practical while preserving reproducible evidence for flagged instances.

6 Measurement Results

We evaluate Evanesca at scale on 14,187,803 Ethereum transactions (2020–2024). The 40 confirmed incidents used in Section 3 span five chains. Section 6.1 details the dataset and collection strategy; Section 6.2 characterizes the pattern landscape at scale (RQ2); Section 6.4 validates flagged defects via source-code inspection (RQ3); and Section 6.6 compares coverage with DeFiRanger.

6.1 Data and Setup

Scope and collection. For January 2020 through December 2024, we collect all Ethereum transactions matching our 40+ adapter set (Section 5), yielding 14,187,803 transactions. For the remaining four chains (BSC, Arbitrum, Avalanche, and Optimism), we retrieve only the transactions associated with our 40 confirmed incidents; these chains contribute RQ1 coverage but not RQ2/RQ3 scale-out. Token values are normalized to USD at each transaction’s block timestamp (Section 5). Instance counts should be interpreted as results over this Ethereum dataset, not as multi-chain population prevalence estimates.

6.2 Pattern Landscape at Scale

We flag 1,418 instances across 14,187,803 transactions using the nine DSL constraints. Table 2 summarizes the distribution across pattern categories.

Table 2. Breakdown of flagged instances at scale (RQ2), including baseline.

Category	Count	Example signal
Price-discrepancy arbitrage (baseline)	956	$ \rho(e) - 1 $ large
Reward-distribution manipulation	207	transient in/out around reward trigger
Price manipulation	166	clustered abnormal swaps
Yield-bearing token deviation	12	persistent share-price gap
Other non-baseline patterns	77	long tail (mixed signals)
Total	1,418	

Baseline separation. Of the 1,418 flagged instances, 956 are routine DEX arbitrage: they exceed $\alpha = 0.05$ but evidence traces show no flash-loan or oracle/lending constraint hit, so we retain them as benign baseline (Figure 3, Panel A). The remaining 462 non-baseline instances combine the trigger with at least one such hit, examined below. We manually validated this partition by cross-checking each instance’s Etherscan event log against its SFG, consulting smart-contract source when action-level evidence was insufficient (Section 6.5).

Non-baseline categories. We group the 462 non-baseline instances by extraction mechanism, yielding three named families plus a long tail. *Price manipulation* (166 instances) uses flash-loan capital to inflate a token’s price on a liquidity pool, then sells previously held tokens at the inflated price on the same pool (Figure 3, Panel B). *Yield-bearing token deviation* (12 instances) reflects gaps between on-chain exchange rates and fair redemption values for derivative tokens. The profit path runs through cross-protocol deposit/withdraw cycles rather than a single corrective swap (Figure 3, Panel C). Section 6.3 gives step-by-step walkthroughs for Panels B and C. *Reward-distribution manipulation* (207 instances) uses flash-loan-funded momentary deposits to capture reward shares without sustained holding (Figure 4). The remaining 77 instances form a long tail that does not cluster into the three named families above.

Finding (RQ2). At 14.2M-transaction scale, the same nine constraints flag 1,418 instances forming a measured pattern landscape: 956 routine arbitrage and 462 non-baseline transactions spanning three distinct families plus a long tail. This resolution shows that DeFi value extraction at scale is structurally diverse, not a single attack template scaled up.




Fig. 3. Representative arbitrage patterns used for the RQ2 classification: routine arbitrage baseline (A), price pumping (B), and yield-bearing token deviation (C).

6.3 Representative Pattern Notes

Figure 3 gives representative evidence traces for the two non-baseline arbitrage families most likely to be confused with routine arbitrage. The notes below summarize the measured action sequence and the constraint signal used for classification.

Panel B: Price Manipulation. The actor borrowed 10.827 ETH from dYdX, used part of the capital to buy 31,997 PENDLE on Sushiswap, and then sold a larger pre-existing PENDLE position into the inflated pool price before repaying the flash loan. The measured net gain was approximately 0.75 ETH in one transaction. The trace fires PRICE_MANIPULATION on the abnormal swap ratio and FLASH_LOAN_ATTACK on the atomic borrow-manipulate-repay cycle.

Panel C: Yield-Bearing Token Deviation. The actor bought Yearn vault-share tokens on Uniswap, deposited them into CreamY, and received Compound cUSDC at a value roughly 2.0% above the fair share price. The cUSDC was then redeemed through Compound and swapped back to ETH, producing approximately 0.097 ETH per cycle. The trace fires EXCHANGE_RATE_MANIPULATION: the cUSDC redeem edge unwraps a position acquired through CreamY at Compound’s fair per-share rate, exposing a cross-protocol exchange-rate gap rather than a simple two-pool arbitrage.

6.4 Previously Unreported Implementation Defects

We inspected all 462 non-baseline instances via event-log and source-code analysis, independently of the extraction-mechanism grouping in Section 6.2. Of the 462 inspected instances, 219 have a harm-causing mechanism we identified in source code, and split into two defect families: all 12 yield-bearing cases map to derivative pricing errors and all 207 reward-manipulation cases to distribution bias. For the remaining 243 (166 price manipulation + 77 long-tail), source-code inspection revealed no implementation defect—extraction is driven by attackers exploiting otherwise correct service logic.² Manual inspection of all 956 baseline instances found no harmful activity (Section 6.5), validating this partition as a reliable triage primitive.

Exchange-rate computation errors. Cross-protocol arbitrageurs captured exchange-rate gaps via deposit/withdraw cycles, unnoticed by protocol operators. Twelve stablecoin-backed derivative tokens (cUSDC and yyDAI families) showed 2–2.5% mispricing per instance, caused by systematic on-chain rate errors where the on-chain rate is computed from the derivative token’s own contract and misses deviations at the same token’s other trading venues. Our detection catches these at $\beta = 0.02$ over lending/oracle-mediated edges (Section 3).

Reward-distribution bias. Several contracts compute payouts based on the balance at the moment of interaction rather than time-weighted holdings, enabling flash-loan-funded deposits that earn disproportionate rewards within a single transaction. Evanescence flags 207 instances, all concentrated in the TEND and WING contracts (shared deflationary logic), cumulatively yielding 15.31 ETH to attackers. Figure 4 illustrates a representative instance (TEND): a flash-loan-funded swap triggers excess reward issuance, and the attacker exits via a reverse swap. The two steps occur in different contracts, requiring cross-protocol evidence traces. Currently both TEND and WING contracts (as well as the yield-bearing wrappers in the previous subsection) are either immutable or no longer actively maintained, limiting the scope for responsible disclosure.

Finding (RQ3). Of the 462 non-baseline instances, 219 (47%) correspond to real implementation defects in two previously uncatalogued families: derivative pricing errors and reward-distribution bias. The same nine constraints thus extend from incident reconstruction (RQ1) to defect discovery, without per-defect tuning.

²We searched DeFiHackLabs, Rekt News, SoK: DeFi Attacks [24], and audit reports for the affected tokens; neither defect family appears in prior aggregated measurement literature.




Fig. 4. Representative reward-distribution manipulation (TEND). A flash-loan-funded swap triggers reward payout and an immediate unwind.

6.5 Flagged Instance Validation and Impact Accounting

Operational measurement yields flagged instances; turning these into actionable findings requires additional validation. We use the following workflow, reproducible from public chain data: (i) re-extract the evidence trace from the transaction identifier and confirm that computed metrics such as $\rho(e)$ and PnL are stable under re-execution; (ii) verify that implied profit exceeds fees/gas and that the value transfer is not a pure accounting artifact; (iii) verify that the operation remains profitable under different reasonable price assumptions, reducing dependence on a single price source.

Impact accounting. We report impact in units native to each extraction path rather than USD totals, avoiding dependence on volatile price sources. For exchange-rate/share-price errors, we express the per-extraction rate gap as a percentage. For reward-distribution manipulation, we compute the attacker’s net profit after all borrowed capital is repaid (e.g., flash-loan repayment); the remaining tokens represent value extracted from the protocol’s reward pool. These figures are conservative lower bounds and should be interpreted as measurement signals grounded in replayable evidence rather than as definitive accounting.

Manual audit of baseline instances. We manually inspected all 956 baseline instances by cross-checking each transaction’s Etherscan event log against Evanescas’s SFG (and smart-contract source when action-level evidence was insufficient). None met our criteria for harmful activity (price manipulation, protocol exploitation, or other adversarial behavior). All instead exhibited cross-protocol price discrepancies (large $|\rho(e) - 1|$ values) with rapid position unwinding consistent with routine arbitrage.

6.6 Comparison with Prior Detection Methods

We compare Evanescas with DeFiRanger, a closely related academic framework for detecting DeFi price manipulation attacks. As it is closed-source, we adopt DeFiScope’s per-incident verdicts [22] for the 12 overlapping incidents and apply DeFiRanger’s published Tr1→Ba→Tr2 pattern for the remaining 28.

Coverage results. Our constraints detect all 40. DeFiRanger covers 12 of 40 (Section 6.6.1). The 23 out-of-scope misses cluster across reentrancy, governance, bridge, oracle defects, admin-key compromise, direct fund injection, and logic-bug mechanisms. The 5 within-scope misses—Pancake Bunny, CreamFinance #2, EGD Finance, Elephant

Money, and Inverse Finance—are PMAs per [22] that DeFiRanger fails on, exposing CFT brittleness across flash-mint chains, share-price defects, and asymmetric (no-Tr2) extractions. Our approach detects all five because constraints evaluate each edge independently.

Within-scope comparison. We restrict the comparison to the 17 incidents within DeFiRanger’s price manipulation scope: DeFiRanger covers 12 (70.6%) and our constraints detect 17 (100%). The gap reflects both coverage breadth and within-scope pattern brittleness. DeFiScope itself misses Harvest #1 (compilation), Indexed Finance (LLM-computation), and Inverse Finance (cross-transaction) among the 12 overlapping incidents [22]. Evanescence catches all three. We can extend Evanescence to a new protocol family by adding only a semantic adapter, without new detection templates (Section 6.6.1).

Architectural difference. DeFiRanger builds cash-flow trees (CFTs) from raw ERC-20 token transfers and pattern-matches Transfer–Balance-change–Transfer (Tr1→Ba→Tr2) sequences, where Tr1 acquires capital, a balance change alters internal state, and Tr2 extracts value at the altered state. Unlike DeFiRanger, our approach constructs SFGs with typed financial actions such as swap, deposit, and borrow, and expresses detection logic as DSL constraints over these edges (Section 4). This enables constraint categories such as reentrancy, governance, bridge, and oracle defects that fixed transfer-level templates cannot express. This explains the 23 out-of-scope misses above: their extraction mechanism is invisible at the transfer level.

6.6.1 Per-Incident Comparison Details. This section provides per-incident outcomes and a constraint-family breakdown for the DeFiRanger comparison summarized in Section 6.6. DeFiRanger is closed-source. We adopt per-incident verdicts from DeFiScope [22] for the 12 incidents in their dataset that overlap with ours, and apply DeFiRanger’s published Tr1→Ba→Tr2 transfer pattern for the remaining 28, classifying each as in-scope detected, in-scope missed, or out-of-scope.

Table 3 reports the per-incident inventory. Evanescence reconstructs every incident in the set, so the table focuses on the DeFiRanger outcome:

- *detected*: the incident’s transfer sequence matches DeFiRanger’s published Tr1→Ba→Tr2 cash-flow-tree pattern.
- *missed*: the incident is a price manipulation, but the published pattern’s structural requirements fail to fire (e.g., bZx’s split capital acquisition breaks the contiguous transfer sequence).
- *out*: outside DeFiRanger’s published price manipulation focus.

Incidents whose value extraction relies on other mechanisms – reentrancy, governance, bridge accounting, share-price oracle defects, admin-key compromise, direct fund injection into pool reserves, or pure logic bugs – are marked *out*. The 23 out-of-scope misses cluster into seven mechanism groups:

- Reentrancy (6): CreamFinance #1, Origin Protocol, Akropolis, Rari Capital, Curve Finance, dForce (the last via read-only reentry on Curve).
- Governance (2): BeanstalkFarms, Fortress Loans (the latter combined with oracle abuse).
- Bridge accounting (2): Qubit Finance, Allbridge.
- Share-price or oracle-misconfiguration defects (3): Yearn Finance, Radiant Capital, Concentric.
- Admin-key compromise (2): Rikkei Finance, Crosswise (oracle overwritten directly).
- Direct fund injection (2): Euler Finance, Hundred Finance (donated funds into pool reserves to inflate share price without a swap step).
- Pure logic bugs (6): KyberSwap, Platypus Finance, BlueberryFDN, Shido Global, Saddle Finance (virtual-price arithmetic), MIM Spell (cauldron precision).

The 5 within-scope misses (Pancake Bunny, CreamFinance #2, EGD Finance, Elephant Money, Inverse Finance) are discussed in Section 6.6.

Table 3. Scope-aware incident inventory used for the RQ1 comparison. Evanescas reconstructs all 40 incidents; the final column records DeFiRanger status.

#	Incident	Loss	Mechanism	Chain	DeFiRanger
2020–2021					
1	CreamFinance #1	\$18.8M	Collat., exch. rate	ETH	out
2	Harvest #1	\$24M	Exch. rate	ETH	detected
3	bZx Hack	\$350K	AMM, collat., oracle, price	ETH	detected
4	Harvest #2	\$10M	Exch. rate	ETH	detected
5	Origin Protocol	\$7M	Collat., exch. rate	ETH	out
6	Value DeFi	\$6M	AMM, price	ETH	detected
7	Warp Finance	\$7.7M	Price	ETH	detected
8	Akropolis	\$2M	Collat., exch. rate	ETH	out
9	Cheese Bank	\$3.3M	AMM, price	ETH	detected
10	Yearn Finance	\$11M	Collat., oracle	ETH	out
11	xToken #2	\$24.5M	AMM, price	ETH	detected
12	Pancake Bunny	\$45M	AMM, price	BSC	missed
13	CreamFinance #2	\$130M	Collat.	ETH	missed
14	Indexed Finance	\$16M	AMM, price	ETH	detected
2022					
15	BeanstalkFarms	\$182M	AMM, price	ETH	out
16	Qubit Finance	\$80M	Bridge	BSC	out
17	Rari Capital	\$80M	Collat., exch. rate	ETH	out
18	EGD Finance	\$36M	AMM, price	BSC	missed
19	Inverse Finance	\$15.6M	AMM, price	ETH	missed
20	Elephant Money	\$11.2M	AMM, price	BSC	missed
21	Float Protocol	\$1.4M	AMM, price	ETH	detected
22	Saddle Finance	\$10M	Collat.	ETH	out
23	Fortress Loans	\$3M	Collat., flash loan	BSC	out
24	Rikkei Finance	\$1.1M	AMM, price	BSC	out
25	Crosswise	\$0.9M	AMM, collat., exch., price	BSC	out
26	Wiener DOGE	\$0.1M	AMM, price	BSC	detected
2023					
27	Euler Finance	\$200M	Collat., exch. rate	ETH	out
28	Curve Finance	\$41M	AMM, price	ETH	out
29	KyberSwap	\$48M	AMM, price	ETH	out
30	Hundred Finance	\$7M	Collat., exch. rate	OP	out
31	Platypus Finance	\$8.5M	AMM, price	AVAX	out
32	Allbridge	\$0.6M	AMM	BSC	out
33	dForce	\$3.7M	Oracle	ARB	out
2024					
34	Radiant Capital	\$4.5M	Oracle	ARB	out
35	WooFi Swap	\$8.8M	AMM, oracle, price	ARB	detected
36	Gamma Strategies	\$3.4M	AMM, oracle	ARB	detected
37	MIM Spell	\$6.5M	AMM, price	ETH	out
38	BlueberryFDN	\$1.4M	Price	ETH	out
39	Concentric	\$1.8M	Oracle	ARB	out
40	Shido Global	\$3.3M	AMM, price	BSC	out
Evanescas overall reconstruction				40/40	
DeFiRanger overall detection				12/40	
DeFiRanger within PMA scope				12/17	

7 Related Work and Discussion

DeFi incident analysis. Prior academic systems analyze DeFi attacks through narrow representations: symbolic oracle analysis with invariant monitoring [18], asset-flow diagrams for individual flash-loan events [2], and attack parameter optimization over DeFi state [15]. More recent frameworks such as DeFort [21] and DeFiScope [22] target price manipulation attacks specifically, tailored to that single attack class. CPMMX [9] automatically synthesizes exploits for constant-product market maker (CPMM) composability bugs, but unlike our retrospective measurement, CPMMX targets exploit generation for a specific AMM family. Unlike these systems, our approach uses invariants that apply to all financial actions such as swaps, deposits, and borrows, regardless of the specific service, enabling coverage of patterns not anticipated by predefined templates.

MEV and operational monitoring. MEV research shows that significant value extraction comes from automated ordering strategies rather than software vulnerabilities [1, 7, 19]. Our patterns overlap but focus on repeatable sequences that exploit protocol mechanics rather than optimizing extraction strategies. Industry monitors such as Forta [8] and OpenZeppelin Defender [13] emphasize real-time alerting but struggle with repeated gray patterns

lacking an obvious exploit trace. Our approach makes such patterns analyzable with explicit rules and replayable evidence that can be re-run across studies.

Key measurement insights. Our results yield three broader observations. First, repeatable DeFi extraction arises from how DeFi primitives such as lending, swapping, and yield generation work, rather than from how any single protocol implements them: the same small constraint set recurs across protocols, forks, and chains (Table 2). Second, our nine constraints target invariants over DeFi primitive actions; extending coverage to orthogonal behaviors, such as gas-cost anomalies or manipulation of supply mechanisms, would require a separate rule family, which we leave to future work. Third, the 309 ms mean processing time (Section 5) permits online circuit-breaking within attack sessions.

Threats to validity. Our confirmed-incident set is curated from public post-mortem reports and may not represent all attack categories. Our semantic adapters cover major AMM, lending, flash-loan, and bridge protocols but not all DeFi primitives (e.g., options, perpetuals). Patterns in unsupported protocols would be missed. The $\alpha = 0.05$ threshold is grounded in prior findings [23], set conservatively above the 1.16% average unexpected slippage on Uniswap reported there. Because confirmed defects span multiple constraint families beyond the swap-ratio check, the RQ3 finding is not dominated by this single parameter. A formal α/β sensitivity surface is deferred to deployment-specific calibration.

Limitations. The nine constraints target invariants over DeFi primitive actions and may miss purely computational vulnerabilities (e.g., access-control bypass) that do not alter observable financial outputs. The 219 confirmed defects (47% of 462) are also a conservative floor; the remaining 243 are price manipulation attacks or long-tail patterns with no identifiable protocol code defect.

8 Conclusion

We presented Evanesca, a measurement pipeline that reconstructs Semantic Financial Graphs from on-chain activity and evaluates reusable DSL constraints over typed financial actions. Across 14,187,803 Ethereum transactions (2020–2024) and 40 confirmed incidents spanning five chains, nine constraints in 107 lines reconstruct every incident, flag 1,418 operational instances spanning five pattern categories, and surface 219 previously unreported defects across two categories, all without per-incident tuning. These results establish Evanesca as a reusable starting point for longitudinal DeFi risk measurement.

Acknowledgments

Open-source release note. The public release includes source code, documentation, reproducible artifact manifests, and example JSON exports for bZx and Harvest replayed transactions. The compiled technical report PDF is intended to document the research prototype and measurement snapshot; it is not a production security guarantee or an automatic attribution system.

References

- [1] Crypto Briefing. 2021. What Is MEV? Ethereum’s Invisible Tax Explained. <https://cryptobriefing.com/what-is-mev-ethereums-invisible-tax-explained/>. Accessed: 2026-04-30.
- [2] Yixin Cao, Chuanwei Zou, and Xianfeng Cheng. 2021. Flashot: A Snapshot of Flash Loan Attack on DeFi Ecosystem. *arXiv preprint arXiv:2102.00626* 1 (2021), 1 – 25.
- [3] ChainLink. [n. d.]. Off-chain price oracle - securely connect smart contracts with off-chain data and services. <https://chain.link/>. Accessed: 2026-04-20.
- [4] CoinGecko. 2024. CoinGecko API: Historical Cryptocurrency Prices and Market Data. <https://www.coingecko.com/en/api>. Accessed: 2026-04-16.
- [5] CoinMarketCap. [n. d.]. CoinMarketCap API: Historical Cryptocurrency Prices and Market Data. <https://coinmarketcap.com/api/documentation/>. Accessed: 2026-04-20.

- [6] Josselin Feist, Gustavo Grieco, and Alex Groce. 2019. Slither: A Static Analysis Framework for Smart Contracts. In *2019 IEEE/ACM 2nd International Workshop on Emerging Trends in Software Engineering for Blockchain (WETSEB)*. IEEE, Montreal, QC, Canada, 8–15.
- [7] Flashbots. 2024. Flashbots: Illuminating the Dark Forest. <https://www.flashbots.net/>. Accessed: 2026-04-06.
- [8] Forta Foundation. 2022. Forta: A Decentralized Runtime Security Solution for Automated Threat Detection and Prevention. <https://docs.forta.network/en/latest/2022-7-11%20Forta%20Litepaper.pdf>. Forta Litepaper. Accessed: 2026-04-06.
- [9] Sujin Han, Jinseo Kim, Sung-Ju Lee, and Insu Yun. 2025. Automated Attack Synthesis for Constant Product Market Makers. *Proceedings of the ACM on Software Engineering 2*, ISSTA (2025), 47–68. doi:10.1145/3728872
- [10] Immunefi. 2025. Crypto Losses in Q1 2025 Report. https://downloads.ctfassets.net/t3wqy70tc3bv/kicQd50NcNkEQyA1QmpOB/b365451e130c1b1c0cbe55c191a22d41/Immunefi_Crypto_Losses_In_Q1_2025.pdf. Accessed: 2026-04-21.
- [11] Bowen Liu, Pawel Szalachowski, and Jianying Zhou. 2021. A First Look into DeFi Oracles. In *2021 IEEE International Conference on Decentralized Applications and Infrastructures (DAPPS)*. IEEE, 39–48.
- [12] Loi Luu, Duc-Hiep Chu, Hrishi Olickel, Prateek Saxena, and Aquinas Hobor. 2016. Making Smart Contracts Smarter. In *Proceedings of the 2016 ACM SIGSAC conference on computer and communications security*. ACM, Vienna Austria, 254–269.
- [13] OpenZeppelin. 2024. OpenZeppelin Defender: Smart Contract Monitoring with Rule-based Sentinels. <https://docs.openzeppelin.com/defender/module/monitor>. Accessed: 2026-04-06.
- [14] Peckshield. 2021. bZx Hack Full Disclosure. <https://peckshield.medium.com/bzx-hack-full-disclosure-with-detailed-profit-analysis-e6b1fa9b18fc>. Accessed: 2026-04-30.
- [15] Kaihua Qin, Liyi Zhou, Benjamin Livshits, and Arthur Gervais. 2021. Attacking the DeFi Ecosystem with Flash Loans for Fun and Profit. In *International Conference on Financial Cryptography and Data Security*. Springer, Springer, Grenada, 3–32.
- [16] Michael Rodler, Wenting Li, Ghassan O Karame, and Lucas Davi. 2019. Sereum: Protecting Existing Smart Contracts Against Re-Entrancy Attacks. In *Proceedings of the Network and Distributed System Security Symposium (NDSS)*. Internet Society, San Diego, CA.
- [17] Petar Tsankov, Andrei Marian Dan, Dana Drachler-Cohen, Arthur Gervais, Florian Bünzli, and Martin T. Vechev. 2018. Securify: Practical Security Analysis of Smart Contracts. In *Proceedings of the 2018 ACM SIGSAC Conference on Computer and Communications Security*. ACM, Toronto, Canada, 67–82.
- [18] Bin Wang, Han Liu, Chao Liu, Zhiqiang Yang, Qian Ren, Huixuan Zheng, and Hong Lei. 2021. BLOCKEYE: Hunting for DeFi Attacks on Blockchain. In *2021 IEEE/ACM 43rd International Conference on Software Engineering: Companion Proceedings (ICSE-Companion)*. IEEE, IEEE, Virtual, 17–20.
- [19] Ben Weintraub, Christof Ferreira Torres, Cristina Nita-Rotaru, and Radu State. 2022. A Flash(bot) in the Pan: Measuring Maximal Extractable Value in Private Pools. In *Proceedings of the 22nd ACM Internet Measurement Conference (IMC)*. ACM, 458–471.
- [20] Siwei Wu, Zhou Yu, Dabao Wang, Yajin Zhou, Lei Wu, Haoyu Wang, and Xingliang Yuan. 2024. DeFiRanger: Detecting DeFi Price Manipulation Attacks. *IEEE Transactions on Dependable and Secure Computing* 21, 4 (2024), 4147–4161.
- [21] Maoyi Xie, Ming Hu, Ziqiao Kong, Cen Zhang, Yebo Feng, Haijun Wang, Yue Xue, Hao Zhang, Ye Liu, and Yang Liu. 2024. DeFort: Automatic Detection and Analysis of Price Manipulation Attacks in DeFi Applications. In *Proceedings of the 33rd ACM SIGSOFT International Symposium on Software Testing and Analysis (ISSTA)*. ACM, Vienna, Austria, 1517–1529.
- [22] Juantao Zhong, Daoyuan Wu, Ye Liu, Maoyi Xie, Yang Liu, Yi Li, and Ning Liu. 2025. DeFiScope: Detecting Various DeFi Price Manipulations with LLM Reasoning. *IEEE Transactions on Dependable and Secure Computing* (2025). Available at <https://ieeexplore.ieee.org/document/11334472>; preprint arXiv:2502.11521.
- [23] Liyi Zhou, Kaihua Qin, Christof Ferreira Torres, Duc Viet Le, and Arthur Gervais. 2021. High-Frequency Trading on Decentralized On-Chain Exchanges. In *2021 IEEE Symposium on Security and Privacy (SP)*. IEEE, 428–445.
- [24] Liyi Zhou, Xihan Xiong, Jens Ernstberger, Stefanos Chaliasos, Zhipeng Wang, Ye Wang, Kaihua Qin, Roger Wattenhofer, Dawn Song, and Arthur Gervais. 2023. SoK: Decentralized Finance (DeFi) Attacks. In *IEEE Symposium on Security and Privacy (SP)*. IEEE, San Francisco, CA, 2444–2461.